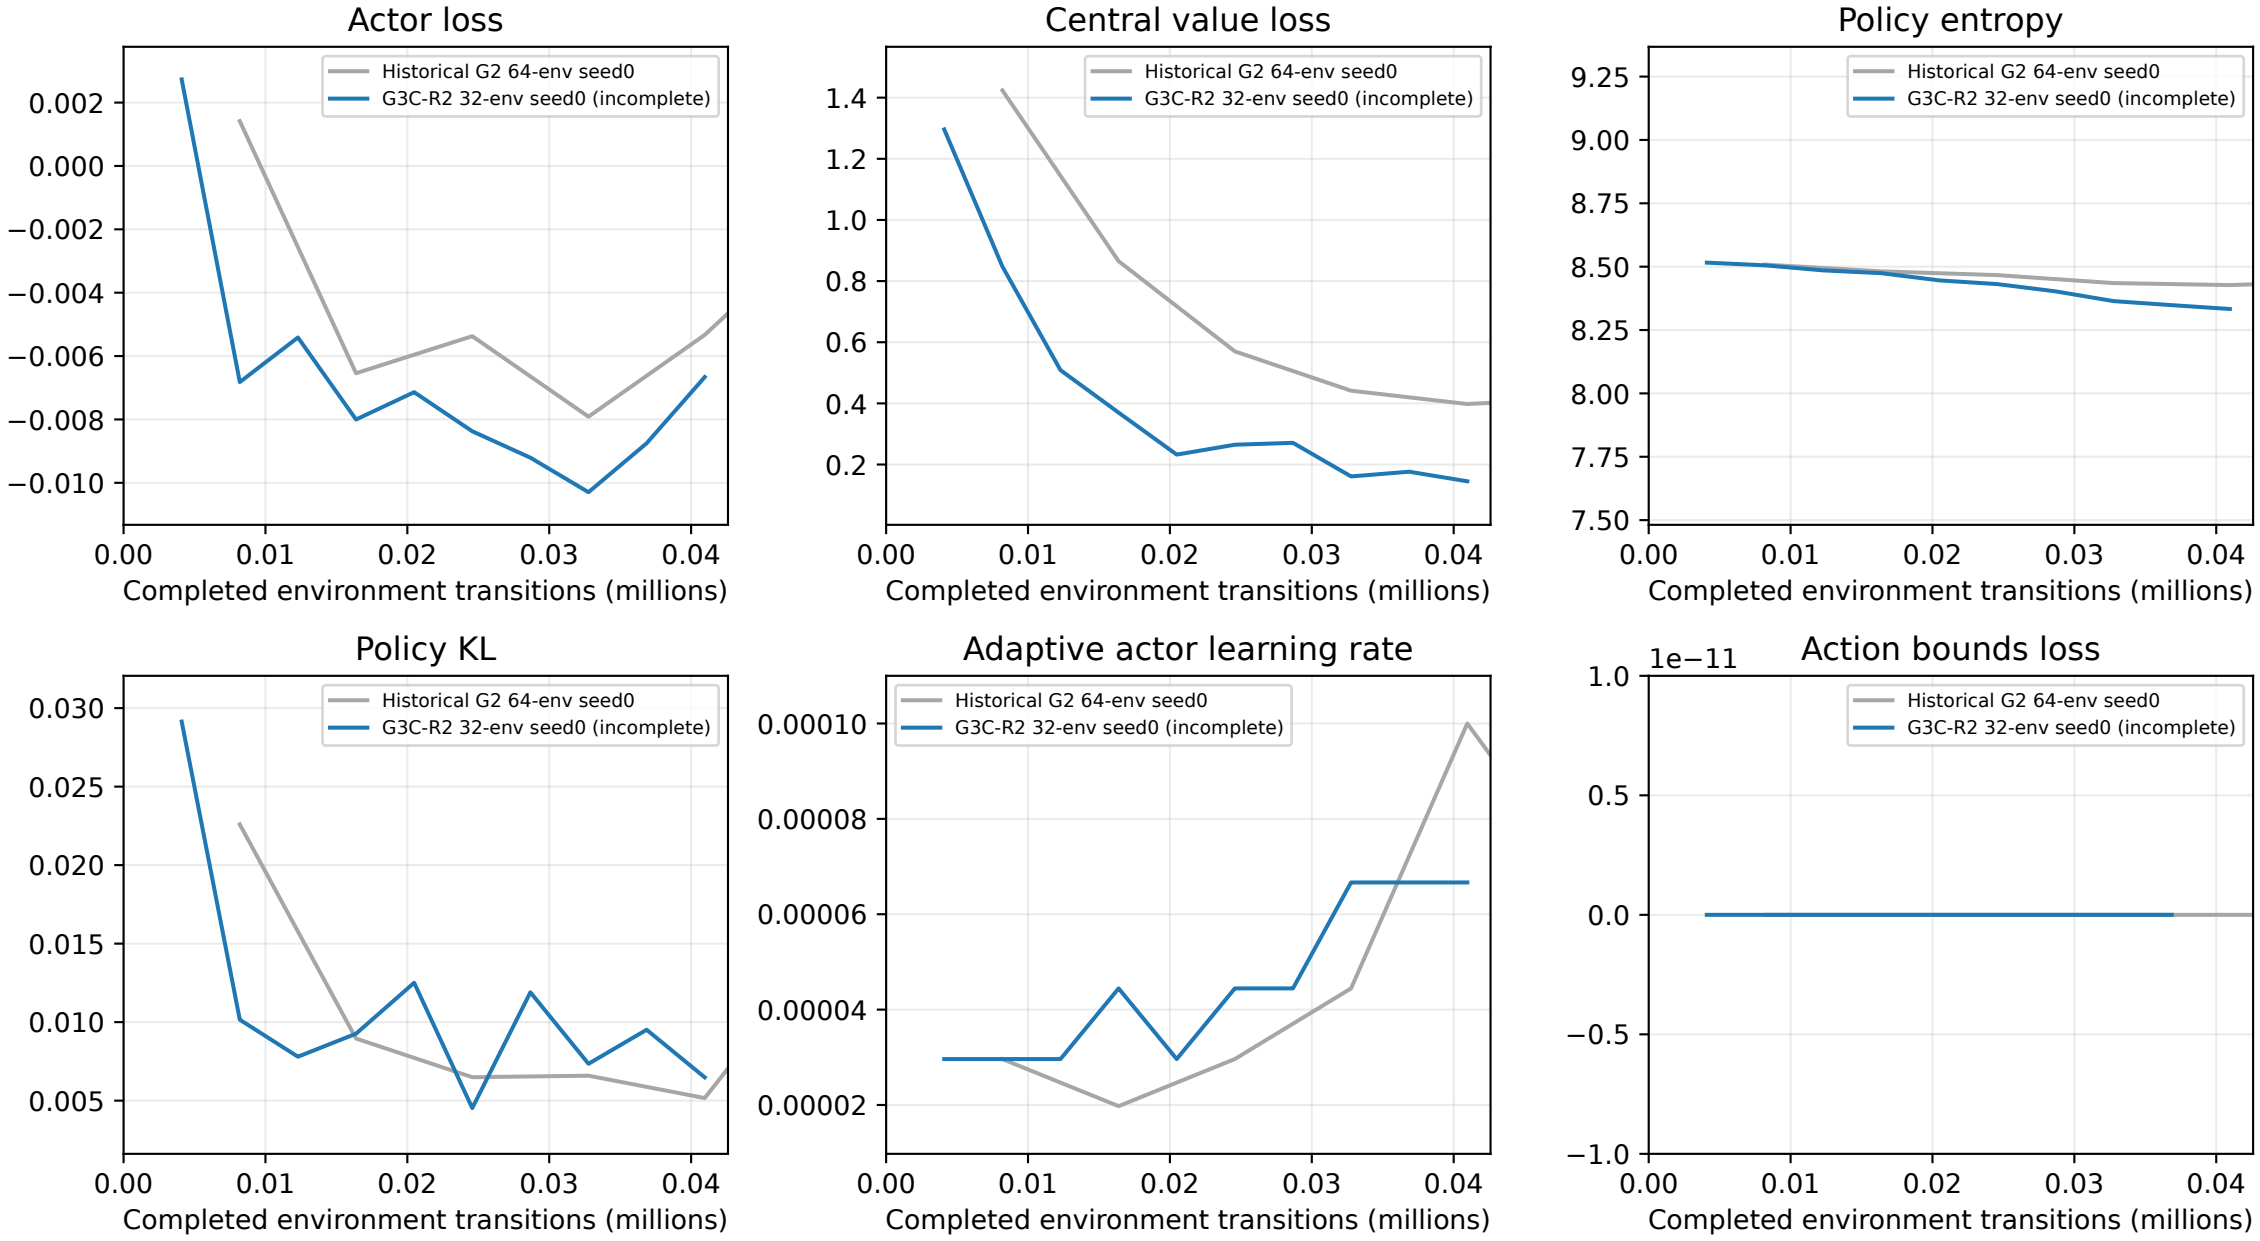
Observed transition prefix; finite logged losses do not prove finite gradients
Actor loss
Central value loss
Policy entropy
Policy KL
Adaptive actor learning rate
Action bounds loss
Historical G2 64-env seed0
G3C-R2 32-env seed0 (incomplete)
Completed environment transitions (millions)
1e−11